## Configuration

**LLM Provider**

- Ollama (local)
- Company API (aivie-exchange)

Base URL

https://aivie-exchange-tnt.sains.c

API Key

••••••••••••••••••••••

Model name

si-gpt-oss-120b

Test

Not yet tested — click Test

TTS Voice

Jenny (US Female)

**Service Status**

Qdrant `localhost:6333`

API not tested

Clear Chat History

# Local Documentation Voice Agent

Drop your docs, ask questions, get voice answers — 100% offline.

## Document Manager

Upload New Documents

Drag & drop files here

Drag and drop files here

Limit 200MB per file • TXT, MD, PDF, HTML, HTM, DOCX

Browse files

Installing Archivematica _ Documentation (Archivema…
421.8KB

1 file(s) selected

Process

**7 document(s) available** — select which to query:

All

None

- [x] **Installing Archivematica _ Documentation (Archivematica 1.18.0) _ Archivematica_ open-source digital preservation system.pdf** NEW
  17 chunks · 2026-03-11 08:35
- [x] **Docbyte Vault local cloud deployments.docx (1).pdf** NEW
  12 chunks · 2026-03-11 08:19
- [x] **Arcsys-requirements-25.2.0.STS-en.pdf**
  255 chunks · 2026-03-10 09:06
- [ ] **DVS_POC_General_Overview.pdf**
  8 chunks · 2026-03-09 15:18
- [ ] **Perbualan dalam Dialek Melayu Sarawak - Lau Annie.mp3_transcript.txt**
  5 chunks · 2026-03-09 14:31
- [ ] **growingplay.com-Did-You-Know-Facts-for-Kids.pdf**
  10 chunks · 2026-03-09 09:54
- [ ] **SAINS-CORP-TP-022 v1.2 - Non-Disclosure Agreement (NDA).docx.pdf**
  19 chunks · 2026-03-09 09:48

🔍 Querying: Arcsys-requirements-25.2.0.STS-en.pdf, Docbyte Vault local cloud deployments.docx (1).pdf, Installing Archivematica _ Documentation (Archivematica 1.18.0) _ Archivematica_ open-source digital preservation system.pdf

## Chat

**what is this**

08:15:41

The passage is simply laying out the key terms that Arcsys uses to describe how it stores and retrieves data. It explains that a lot is a collection of one or more archived objects, and lot enrichment means adding new objects to an existing lot. The manifest is an XML file that acts like an ID card for a record, listing all of the metadata, file names, sizes, hashes and other details about the objects it describes. A metadata element definition (or keyword) defines the type and format of a piece of metadata, and an object is the basic unit that can be archived—a file, directory, database table, etc. In short, the excerpt defines the vocabulary around Arcsys's archiving system.

Play Audio

Sources: `Arcsys-requirements-25.2.0.STS-en.pdf`

**If i were to deploy this Arcsys what server spec i need?**

08:16:38

To run Arcsys you'll need a fairly standard server-class box. At a minimum the software will run on a machine with 4 GB of RAM, but for anything beyond a tiny test deployment you'll want more memory.

For the CPU the guide recommends a recent-generation Intel Xeon and suggests somewhere between eight and thirty-two cores as a comfortable range. The platform is built to scale with additional cores, so if you expect a heavy load you can size up within that band.

Networking-wise the absolute minimum is a 1 Gb Ethernet adapter, but if you're archiving large volumes of data the recommendation bumps up to a 10 Gb Ethernet link for optimal throughput.

On the operating-system side Arcsys is supported on Windows Server 2022 64-bit. If you prefer Linux you can use any modern distro that provides systemd (even if you're not running Arcsys as a systemd service). On Windows you also have to install the Microsoft Visual C++ 2019 Redistributable (x64) and the .NET Framework 4.0 runtime.

Finally, because Arcsys is Java-based you must have a 64-bit Java 17 JRE installed. Both Oracle Java and OpenJDK are supported; some distributions bundle the JRE with the JDK, so just make sure a compatible Java 17 runtime is present.

So, in short: a server with at least 4 GB RAM, an 8-to-32-core recent Xeon CPU, 1 Gb Ethernet (10 Gb recommended for heavy workloads), Windows Server 2022 64-bit (or a Linux distro with systemd), Visual C++ 2019 Redistributable, .NET 4.0, and a 64-bit Java 17 runtime. That will meet the documented prerequisites for an Arcsys deployment.

Play Audio

Sources: `Arcsys-requirements-25.2.0.STS-en.pdf`

**Can i use Cloud for this?**

08:17:21

Yes, you can use cloud archiving with Arcsys, but only if you meet the listed prerequisites. The transfer server must run on a Linux version that Arcsys Core supports, and the system needs libcurl version 4.5 or newer compiled with a supported SSL library (you can check the compatible SSL libraries on the libcurl website). Your cloud-storage endpoint also has to speak the Amazon S3 REST API – things like HEAD requests for bucket and object existence, POST …?uploads for multipart starts and PUT …?partNumber=… for each part. In practice that means using a MinIO or Ceph Object Gateway of version 10.1 or later, or any other S3-compatible service that implements those calls. Object Lock is technically compatible, though Arcsys won't configure it for you. So as long as those requirements are satisfied, cloud archiving will work.

Play Audio

Sources: `Arcsys-requirements-25.2.0.STS-en.pdf`

**What happen if Vendor need to use their country's Law instead of ours ?**

08:18:17

The NDA explicitly says that any questions about how the agreement works – its validity, how it's interpreted and enforced – are governed by Malaysian law. That means the vendor can't simply switch to the laws of their own country unless both parties agree to rewrite that clause. If they tried to rely on their own jurisdiction without an amendment, the contract would still be interpreted under Malaysian law, and insisting on a different legal framework could be considered a breach of the agreement. In short, the vendor would need to accept Malaysian law as written, or you'd have to negotiate a change to the applicable-law provision.

Play Audio

Sources: `SAINS-CORP-TP-022 v1.2 - Non-Disclosure Agreement (NDA).docx.pdf`

**What are the consequences? **

08:19:00

If you breach the NDA, the other side can take several actions. First, they may go to court to get an injunction or other equitable relief because money alone won't make up the damage. They can also, at their sole discretion, end any ongoing discussions with you and bar you from taking part in any future tenders or procurement exercises with them, and you won't be entitled to any compensation for that. In addition, you'll be required to return all confidential material – documents, drawings, models, notes, and any copies – promptly and to destroy any remaining copies. The agreement also makes clear that any attempt to claim damages for the breach is unlikely to succeed if the non-breaching party can prove irreparable harm. So the practical consequences are legal injunctions, loss of business opportunities, and the obligation to hand back and destroy all confidential information.

Play Audio

Sources: `SAINS-CORP-TP-022 v1.2 - Non-Disclosure Agreement (NDA).docx.pdf`

**How to deploy Docbyte? What I need to prepare? The specification needed?**

08:20:51

To get Docbyte Vault up and running you'll need a Kubernetes-or-OpenShift cluster that you either share with other workloads or spin up just for Vault. The first thing to do is make sure the cluster can provide the compute, memory and persistent storage that the platform's containers expect.

Each core service runs in its own container, so you'll be deploying a set of pods—things like the ingest service, preservation service, access service, the static UI for Explorer and Review, and an authentication server (Keycloak). All of them talk over secure internal APIs and can be scaled out independently, so you'll want to reserve a bit of extra capacity for horizontal scaling and high availability.

In terms of concrete numbers, the baseline sizing the document gives is:

- Vault Explorer static UI: two pods, each using half a vCPU and 256 MB of RAM, for a total of about 0.5 vCPU and 0.5 GB of RAM.
- Vault Review static UI: the same – two pods, 0.5 vCPU and 256 MB per pod.
- Keycloak (the OpenID Connect or SAML identity provider): the doc doesn't list the exact resources, but you should plan for a similar small footprint and then size it higher if you expect many concurrent users.

Add whatever additional pods you need for ingest, preservation, access, and any other optional components, and multiply the per-pod vCPU and memory by the number of replicas you plan to run. Remember that the figures above are for steady-state operation and don't include the overhead of the cluster itself, monitoring agents, or side-car containers.

You'll also need a MongoDB deployment for metadata and configuration storage. You can either install MongoDB as part of the Vault installation or point Vault at an existing managed MongoDB cluster that supports replica sets and sharding. Using a managed cluster can help you stay in line with internal database governance policies and reduce duplicate infrastructure.

For storage, Docbyte Vault expects object-based storage that speaks the S3 protocol. Make sure you have an S3-compatible bucket (or a service like MinIO, AWS S3, Azure Blob with S3 compatibility, etc.) that the containers can reach.

Other preparation steps include:

- Setting up the OpenID Connect or SAML identity provider (Keycloak) and configuring Vault to use it for authentication.
- Ensuring your logging, monitoring and observability stack is ready so Vault can ship its logs and metrics into it.
- Deciding on your deployment method—most customers use Helm charts or OpenShift Operators, which let you define secrets, config maps and other parameters in a way that fits your CI/CD pipelines.
- Planning for high availability: because every component is containerized, you can run each service in a replica set and use Kubernetes/Openshift built-in features like pod anti-affinity and node pools to keep them redundant.

Once the cluster is provisioned with the right resources, you can apply the provided Helm charts or Operator descriptors, supply the required secrets (like S3 credentials, MongoDB connection strings and Keycloak client details), and let the platform roll out the services. After that, you'll have a horizontally scalable, container-based Docbyte Vault ready for ingest, preservation and access of your digital assets.

Play Audio

Sources: `Docbyte Vault local cloud deployments.docx (1).pdf`

**Summarize specification i need to prepare to deploy this?**

08:24:14

To run Docbyte Vault you'll need a Kubernetes or OpenShift cluster that can supply roughly 48 virtual CPUs and about 129 gigabytes of RAM in total. The baseline topology consists of 23 pods spread across the various services:

- Two static-UI pods for Vault Explorer and Vault Review, each with half a vCPU and 256 MB of memory.
- Keycloak for identity and access, using two pods, each with 2 vCPU and 2 GB of RAM.
- PostgreSQL for configuration and IAM, two pods, each with 1 vCPU and 2 GB of RAM (total 2 vCPU, 4 GB).
- The ingest & dissemination API (or a combined API container for small installs) runs two pods, each with 2 vCPU and 4 GB of RAM.
- Vault Core, the data-management engine, uses two pods, each with 4 vCPU and 16 GB of RAM.
- Kafka for streaming and queuing, three pods, each with 2 vCPU and 8 GB of RAM.
- MongoDB for metadata, three pods, each with 2 vCPU and 8 GB of RAM. Because MongoDB is stateful, its pods should be spread across at least three distinct nodes for resilience.
- OpenSearch for indexing and search, three pods, each with 4 vCPU and 8 GB of RAM, also spread over three nodes.
- Antivirus engine, two pods, each with 2 vCPU and 4 GB of RAM.

Add up the resources per component and you get the 48 vCPU-129 GB-RAM baseline. In addition to compute and memory, each service that stores data—PostgreSQL, MongoDB, OpenSearch, and any Vault Core persistence—needs reliable persistent storage, typically provisioned via the cluster's storage class. The solution expects object-storage access over the S3 protocol for vault-level storage, so you'll also need an S3-compatible bucket (or compatible cloud storage) with enough capacity for your expected ingest volume.

Because the numbers above are for steady-state operation, you should also reserve extra capacity for system overhead, monitoring agents, side-car containers, and future scaling. If your ingest rates, parallel streams, or object sizes are higher than average, you'll want to contact Docbyte for a customized sizing.

In short, plan for a cluster that can run at least 23 pods, delivering 48 vCPU and 129 GB of RAM, with persistent volumes for the databases and an S3-compatible backend, and distribute the MongoDB and OpenSearch pods across three separate nodes to meet the recommended redundancy.

Play Audio

Sources: `Docbyte Vault local cloud deployments.docx (1).pdf`

**how to install archivematica? what spec needed?**

08:35:49

To get Archivematica up and running you'll first need a supported Linux box—Ubuntu 24.04 is the reference platform. Make sure you have root or sudo access, then follow these high-level steps:

1. Update your system and install the basic packages the software expects, such as Python 3, Git, curl and a recent Java runtime. On Ubuntu the default OpenJDK 11 works, but the Archivematica components that talk to Elasticsearch still need Java 8, so you'll want to add the OpenJDK 8 package as well.

2. Install a supported relational database. Archivematica 1.18 has been tested with MySQL 8.0, including Percona or MariaDB drops that are compatible with MySQL 8.0. Create a dedicated database and user for Archivematica, then grant that user full rights on the new database.

3. (If you want the search features) install Elasticsearch. Since version 1.7 the search engine is optional, but most installations keep it. The documentation notes that OpenJDK 8 is the Java version that ships with Elasticsearch on Ubuntu 24.04, so make sure that's present.

4. Pull the Archivematica packages. The easiest way on Ubuntu is to add the Archivematica APT repository, update the package list, and then install the "archivematica-pipeline" and "archivematica-storage-service" packages. The installer will also bring in the exact versions of the remaining Python libraries and system utilities that Archivematica expects, so you don't need to tweak those.

5. Run the automated configuration scripts. After the packages are in place you'll run the archivematica-setup wizard, which will ask you for the database credentials, the hostname of the storage service, and whether you want SSL enabled. If you choose SSL, the wizard will generate self-signed certificates for you, which you can replace later with your own.

6. Start the services. The setup creates systemd units for the pipeline, the storage service, and for Elasticsearch if you installed it. Enable and start them with systemctl enable and systemctl start. You can verify everything is running by visiting the web interfaces: http://your-host:8000 for the pipeline UI and http://your-host:8001 for the storage service UI.

7. Do a quick test. Upload a small bag or a simple directory through the pipeline UI, run the default workflow, and confirm the object appears in the Storage Service and that the search index updates. If everything looks good you're ready to start configuring your own processing workflows.

**Hardware specifications**

The documentation lists two tiers:

*Minimum hardware* – a single virtual machine with at least two CPU cores, 4 GB of RAM and 80 GB of disk space. That's enough for a basic test or a very small pilot.

*Recommended production* – a dual-core or better server with 8 GB or more RAM, 200 GB of fast storage (SSD is ideal) and a separate host for Elasticsearch if you expect heavier indexing loads. For larger collections you'll want to scale out: add more CPU and RAM, spread the pipeline workers across additional machines, and consider a dedicated Elasticsearch cluster.

Finally, make sure you're using a modern web browser—Chrome, Firefox or Edge are all supported—for the user interfaces.

That covers the essential steps and the hardware you'll need to get Archivematica installed and running. If you hit any snags, the "Advanced installation" section of the docs goes into SSL configuration, disabling task-output capture, and scaling the pipeline across multiple machines. Good luck!

Play Audio

Sources: Installing Archivematica _ Documentation (Archivematica 1.18.0) _ Archivematica_ open-source digital preservation system.pdf

Ask a question about the selected documents...

**Which is the easiest to be deploy and setup ? Can you rank them from easiest to hardest? and explain why**

If you're looking for the path of least resistance, the Docbyte Vault Helm-chart deployment is the simplest. It's basically a single helm install that pulls in all the containers, wires up the secrets and config, and leaves you with a running system in a reproducible way.

Next up are the OpenShift Operators for Docbyte Vault. They're a bit more involved because you first have to have an OpenShift cluster and then create the custom resource that the operator watches, but once that's done the operator handles upgrades, scaling and health checks for you, so it's still fairly hands-off.

On the Archivematica side the easiest automated route is the provided Ansible playbooks. You run the playbook against a fresh Rocky Linux or Red Hat host, it installs the packages, configures the services and even sets up the processing-disk recommendations for you. It does require a bit of Ansible know-how – inventory files, variables, maybe a tweak for SSL if you need it – but it's a single, repeatable command.

Finally, the most manual option is installing Archivematica directly from the distribution packages. You have to resolve dependencies, edit configuration files, set up the database, configure the storage backend and handle SSL yourself. It works, but you'll spend the most time in the terminal figuring out each step.

So, ranking from easiest to hardest:

1. Docbyte Vault via a Helm chart – one-line install, everything templated.
2. Docbyte Vault via an OpenShift Operator – requires an OpenShift cluster but the operator manages the lifecycle.
3. Archivematica via the Ansible scripts – automated but needs some Ansible setup.
4. Archivematica via manual package installation – most hands-on, many individual configuration steps.

The order reflects how much of the heavy lifting is already done for you versus how many manual decisions and commands you have to make.

Play Audio

0:01 / 1:59

Download

Sources: Docbyte Vault local cloud deployments.docx (1).pdf, Installing Archivematica _ Documentation (Archivematica 1.18.0) _ Archivematica_ open-source digital preservation system.pdf